

Where things stand

The full build pipeline, what's done, what's next, and everything intentionally parked for later — as of 24 Aug 2026.

5

MASTER-DATA MODULES DONE

2

WORKFORCE ITEMS
PARTIALLY STARTED

1

UP NEXT

9

PIPELINE STAGES NOT
STARTED

14

DECISIONS PARKED ON
PURPOSE

The build pipeline

Priority order for this project: Usability → Operational Correctness → Data Integrity → Scalability → Advanced Features. The stages below are built roughly in this sequence — each one leans on the ones above it.

Master Data

● 5 of 6 done

The foundation everything else is built on. Auth/multi-tenancy plus four full masters are live; Locations/Bins is the one piece left.

— Auth + Multi-Tenancy

JWT login, tenant isolation proven across two real test companies, zero cross-tenant leakage.

● Done

— Warehouse Master

Doubles as a network-node master (Factory/Distributor/DC/CNF/Cross-dock), storage-type & dispatch-flow breakdowns, bulk import, Delete All.

● Done

— SKU Master

Full CRUD, category as a real FK, bulk import/export, summary analytics, Delete All.

● Done

— Customer Master

Multi ship-to per customer, Local/Upcountry delivery-zone tagging, bulk import, Delete All.

● Done

— User Master

Hierarchical roles & warehouse scoping, bulk import, self-service edit, login ledger + history viewer. Built this session.

● Done

— Locations / Bins

Schema exists, nothing built on top yet. Needs its own design pass first: numbering scheme, aisle/rack/level structure, how pallet counts get carved out by zone.

● Up next

2

Yard & Gate Management

● Not started

What happens before a truck ever reaches a dock — added to the pipeline this round, sits ahead of Inbound in the real physical flow.

Gate In / Gate Out

Yard Management

Dock Scheduling (In & Out)

Dock Door Management

Weight Bridge Integration

3

Inbound (Receipts)

● Not started

Recording goods arriving at a warehouse. Depends on Locations/Bins existing to receive into.

ASN Receipts & Tracker

QC

Cross-Dock

4

Putaway

● Not started

Moving received stock from dock to its storage bin — the flow point exists in the build order, the actual putaway logic (rules for where a SKU lands) is what's still missing.

5

Inventory (ledger + live views)

● Not started

The core invariant is already locked in at the schema level: `StockMovement` is an append-only ledger — every receipt, putaway, pick, dispatch, return, and adjustment inserts a signed-quantity row, on-hand stock is always derived, never stored directly. The live views/screens on top of it aren't built.

Cycle Count & Consolidation

Re-Slotting

SKU Categorization (velocity/ABC)

6

Outbound (Orders / Allocation)

● Not started

Customer orders and allocating stock against them.

7

Picking

● Not started

Turning an allocated order into a floor task.

[Single / Multi Order Pick](#)[Wave Picking & Other Picking Logics](#)

8

Dispatch

● Not started

Loading and sending stock out. Dispatch-proximity distance calc and cutoff-time policy are parked decisions that belong here once this stage starts (see below).

[Truck Builder](#)

9

Returns

● Not started

Stock coming back in.

10

Analytics

● Not started

Deliberately last — built on top of everything above, including the Workforce & Labor items below.

[Dashboard for All KPIs](#)

Workforce & Labor Management

Cross-cutting, not a single sequential stage — these run alongside the whole pipeline rather than after one particular stage. User Master (Master Data, above) is the seed of this: role/warehouse structure and the login ledger both already exist.

Workforce Management

● Partially started

User Master's role hierarchy and warehouse scoping is the structural seed — assigning who works where, and who can see whom.

Attendance Management

● Partially started

The login ledger (LoginEvent) already records every login permanently, and the User Master list shows Days Active/Last Login. The actual daily-attendance report/rollup built on top of that raw data isn't built.

Roster Planning

● Not started

Shift scheduling — who's meant to be working when.

Productivity Analysis

● Not started

Per-person/team output — ties into Analytics once there's transactional data (picks, putaways) to measure against.

Parked on purpose

Not forgotten — each of these was raised and consciously set aside, grouped by what kind of decision it is. None of these are pipeline stages; they can slot in whenever the moment's right, independent of the sequence above.

Access & Security

- **Super Admin creation** — still no controlled way to create the platform-level role.
- **Configurable permission matrix** — today's role rules are fixed, not client-editable.
- **Network-restricted logins** — locking a shop-floor ID login to the client's own network.
- **Invite-link flow** — accounts get a password set directly today, no emailed invite.

Data & Policy

- **SkuRelationship** — kits/combos, schema exists, no logic behind it.
- **CategoryPackSpec data** — seeded empty, no real rows yet.
- **Inventory Control Policy** — min/max, reorder point, FIFO/FEFO/LIFO.
- **Opening Balance load** — bringing existing stock into the ledger at go-live.
- **Dispatch-proximity calc** — lat/long fields exist, no distance algorithm.
- **Dispatch cutoff-time policy** — belongs to a future Dispatch Policy stage.

Commercial

- **3PL billing / contract** — billing a client based on storage/handling terms.

Infrastructure

- **Real SAP/ERP integration** — `expCode` fields exist as a landing spot, nothing behind them.
- **TMS integration** — Transportation Management System, carrier/route/freight.
- **Cloud/production deployment** — local Docker Compose (Postgres) only, so far.

Reflects CLAUDE.md as of commit 5c178ca (24 Aug 2026), plus the fuller feature scope added the same day. This is a growing list, not a fixed one – bring more items whenever they come to mind and this page gets extended, not replaced. A snapshot, not a live document – regenerate whenever status changes meaningfully.